

decoding-default

August 4, 2025

##Mount Drive

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

##Import Required Libraries

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
import matplotlib.pyplot as plt
```

##Install Data profiling

```
[ ]: !pip install ydata-profiling
```

##Load Dataset

```
[ ]: df=pd.read_csv('/content/drive/MyDrive/Lending_Data/lending_club_loan_two.csv')
```

##Data Profiling & Summary

```
[ ]: # data profiling - Run only once for the first time
#from ydata_profiling import ProfileReport

#profile = ProfileReport(df, title="Data Profiling Report", explorative=True)

#profile.to_file("output.html")
```

```
[ ]: # Basic Info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---
```

```

0   loan_amnt          396030 non-null float64
1   term               396030 non-null object
2   int_rate          396030 non-null float64
3   installment       396030 non-null float64
4   grade             396030 non-null object
5   sub_grade         396030 non-null object
6   emp_title         373103 non-null object
7   emp_length        377729 non-null object
8   home_ownership    396030 non-null object
9   annual_inc        396030 non-null float64
10  verification_status 396030 non-null object
11  issue_d           396030 non-null object
12  loan_status       396030 non-null object
13  purpose           396030 non-null object
14  title             394274 non-null object
15  dti               396030 non-null float64
16  earliest_cr_line  396030 non-null object
17  open_acc          396030 non-null float64
18  pub_rec           396030 non-null float64
19  revol_bal         396030 non-null float64
20  revol_util        395754 non-null float64
21  total_acc         396030 non-null float64
22  initial_list_status 396030 non-null object
23  application_type  396030 non-null object
24  mort_acc          358235 non-null float64
25  pub_rec_bankruptcies 395495 non-null float64
26  address           396030 non-null object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB

```

```
[ ]: df.describe()
```

```

[ ]:
count      loan_amnt      int_rate      installment      annual_inc  \
count  396030.000000  396030.000000  396030.000000  3.960300e+05
mean    14113.888089    13.639400    431.849698  7.420318e+04
std      8357.441341     4.472157    250.727790  6.163762e+04
min       500.000000     5.320000     16.080000  0.000000e+00
25%      8000.000000    10.490000    250.330000  4.500000e+04
50%     12000.000000    13.330000    375.430000  6.400000e+04
75%     20000.000000    16.490000    567.300000  9.000000e+04
max     40000.000000    30.990000   1533.810000  8.706582e+06

count      dti      open_acc      pub_rec      revol_bal  \
count  396030.000000  396030.000000  396030.000000  3.960300e+05
mean     17.379514    11.311153     0.178191  1.584454e+04
std      18.019092     5.137649     0.530671  2.059184e+04
min        0.000000     0.000000     0.000000  0.000000e+00

```

25%	11.280000	8.000000	0.000000	6.025000e+03
50%	16.910000	10.000000	0.000000	1.118100e+04
75%	22.980000	14.000000	0.000000	1.962000e+04
max	9999.000000	90.000000	86.000000	1.743266e+06

	revol_util	total_acc	mort_acc	pub_rec_bankruptcies
count	395754.000000	396030.000000	358235.000000	395495.000000
mean	53.791749	25.414744	1.813991	0.121648
std	24.452193	11.886991	2.147930	0.356174
min	0.000000	2.000000	0.000000	0.000000
25%	35.800000	17.000000	0.000000	0.000000
50%	54.800000	24.000000	1.000000	0.000000
75%	72.900000	32.000000	3.000000	0.000000
max	892.300000	151.000000	34.000000	8.000000

```
[ ]: df.head()
```

```
[ ]:   loan_amnt      term  int_rate  installment  grade  sub_grade  \
0    10000.0    36 months    11.44         329.48     B         B4
1     8000.0    36 months    11.99         265.68     B         B5
2    15600.0    36 months    10.49         506.97     B         B3
3     7200.0    36 months     6.49         220.65     A         A2
4    24375.0    60 months    17.27         609.33     C         C5
```

```
      emp_title  emp_length  home_ownership  annual_inc  ...  \
0      Marketing    10+ years          RENT    117000.0  ...
1  Credit analyst     4 years    MORTGAGE    65000.0  ...
2  Statistician    < 1 year          RENT    43057.0  ...
3  Client Advocate     6 years          RENT    54000.0  ...
4  Destiny Management Inc.    9 years    MORTGAGE    55000.0  ...
```

```
open_acc  pub_rec  revol_bal  revol_util  total_acc  initial_list_status  \
0      16.0      0.0    36369.0         41.8        25.0                w
1      17.0      0.0    20131.0         53.3        27.0                f
2      13.0      0.0    11987.0         92.2        26.0                f
3       6.0      0.0     5472.0         21.5        13.0                f
4      13.0      0.0    24584.0         69.8        43.0                f
```

```
application_type  mort_acc  pub_rec_bankruptcies  \
0      INDIVIDUAL        0.0                0.0
1      INDIVIDUAL        3.0                0.0
2      INDIVIDUAL        0.0                0.0
3      INDIVIDUAL        0.0                0.0
4      INDIVIDUAL        1.0                0.0
```

```
address
0    0174 Michelle Gateway\r\nMendozaberg, OK 22690
```

```

1  1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113
2  87025 Mark Dale Apt. 269\r\nNew Sabrina, WV 05113
3      823 Reid Ford\r\nDelacruzside, MA 00813
4      679 Luna Roads\r\nGreggshire, VA 11650

```

[5 rows x 27 columns]

##Data Cleaning

```
[ ]: df.shape
```

```
[ ]: (396030, 27)
```

###Removing Missing Values

```
[ ]: #Remove missing values from features emp_length, mort_acc, pub_rec_bankruptcies
df.dropna(subset=['emp_length', 'mort_acc', 'pub_rec_bankruptcies'],
           inplace=True)
df.shape
```

```
[ ]: (340990, 27)
```

###Removing Zeroes

```
[ ]: #Remove zeroes from features such as annual_inc, revol_bal, revol_util
df = df[df['annual_inc'] != 0]
df = df[df['revol_bal'] != 0]
df = df[df['revol_util'] != 0]
df.shape
```

```
[ ]: (339541, 27)
```

```
[ ]: #Remove none and other values from the rows of home_ownership
df = df[df['home_ownership'] != 'NONE']
df = df[df['home_ownership'] != 'OTHER']
df.shape
```

```
[ ]: (339482, 27)
```

###Plot Outliers






```
[ ]: #Plot outliers
# Select numerical columns
numeric_cols = df.select_dtypes(include='number').columns






# Plot boxplots for each numeric column
for col in numeric_cols:
    plt.figure(figsize=(8, 1.5))
```



```

sns.boxplot(x=df[col], color='orange')
plt.title(f"Outlier Detection - {col}")
plt.xlabel(col)
plt.tight_layout()
plt.show()

```







###Removing Outliers

```
[ ]: #Remove outliers from features loan_amnt, int_rate, open_acc, revol_util
# List of target columns to remove outliers from
cols_to_clean = ['loan_amnt', 'int_rate', 'open_acc', 'revol_util']

# Loop through each column and remove outliers using IQR method
for col in cols_to_clean:
    if col in df.columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        # Filter out outliers
        df = df[(df[col] >= lower_bound) & (df[col] < upper_bound)]

print(f"Outliers removed. New shape of DataFrame: {df.shape}")
```

Outliers removed. New shape of DataFrame: (325008, 27)

```
[ ]: #Remove outliers and change upper limit of :
# Remove rows where annual_inc > 20 lakhs
if 'annual_inc' in df.columns:
    df = df[df['annual_inc'] <= 2_000_000]
#Remove rows from total_acc > 110
if 'total_acc' in df.columns:
    df = df[df['total_acc'] <= 110]
#Remove rows from revol_bal > 7.5 lakhs
```

```

if 'revol_bal' in df.columns:
    df = df[df['revol_bal'] <= 7_500_000]
#Remove rows from mort_acc >15
if 'mort_acc' in df.columns:
    df = df[df['mort_acc'] <= 15]
    #Remove rows from pub_rec >15
if 'pub_rec' in df.columns:
    df = df[df['pub_rec'] <= 15]
    #Remove rows from pub_rec_bankruptcies >15
if 'pub_rec_bankruptcies' in df.columns:
    df = df[df['pub_rec_bankruptcies'] <= 15]
#Remove rows from dti >70
if 'dti' in df.columns:
    df = df[df['dti'] <= 70]
    #Remove rows from installment > 1450
if 'installment' in df.columns:
    df = df[df['installment'] <= 1450]

```



```
[ ]: print(f"Outliers removed. New shape of DataFrame: {df.shape}")
```






Outliers removed. New shape of DataFrame: (324859, 27)






```

[ ]: # Plot boxplots for each numeric column after outlier removal
for col in numeric_cols:
    plt.figure(figsize=(8, 1.5))
    sns.boxplot(x=df[col], color='orange')
    plt.title(f"After Outlier Removal - {col}")
    plt.xlabel(col)
    plt.tight_layout()
    plt.show()

```







```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 324859 entries, 0 to 396028
```

Data columns (total 27 columns):

#	Column	Non-Null Count	Dtype
0	loan_amnt	324859 non-null	float64
1	term	324859 non-null	object
2	int_rate	324859 non-null	float64
3	installment	324859 non-null	float64
4	grade	324859 non-null	object
5	sub_grade	324859 non-null	object
6	emp_title	321562 non-null	object
7	emp_length	324859 non-null	object
8	home_ownership	324859 non-null	object
9	annual_inc	324859 non-null	float64
10	verification_status	324859 non-null	object
11	issue_d	324859 non-null	object
12	loan_status	324859 non-null	object
13	purpose	324859 non-null	object
14	title	323457 non-null	object
15	dti	324859 non-null	float64
16	earliest_cr_line	324859 non-null	object
17	open_acc	324859 non-null	float64
18	pub_rec	324859 non-null	float64
19	revol_bal	324859 non-null	float64
20	revol_util	324859 non-null	float64
21	total_acc	324859 non-null	float64
22	initial_list_status	324859 non-null	object
23	application_type	324859 non-null	object
24	mort_acc	324859 non-null	float64
25	pub_rec_bankruptcies	324859 non-null	float64
26	address	324859 non-null	object

dtypes: float64(12), object(15)

memory usage: 69.4+ MB

#Feature Engineering

```
[ ]: #Feature Engineering for address column - extract zipcode from address
# Extract 2-letter state code and 5-digit pincode from address
df['zip_code'] = df['address'].str.extract(r'([A-Z]{2}\s*\d{5})$')
```

##Extract necessary data Conversion to date time format

```
[ ]: from datetime import datetime
```

```
[ ]: # Step 1: Convert to datetime (if not already)
df['issue_d'] = pd.to_datetime(df['issue_d'], errors='coerce')
df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'], errors='coerce')

# Step 2: Strip the time (keep only date)
```

```
df['issue_d'] = df['issue_d'].dt.date
df['earliest_cr_line'] = df['earliest_cr_line'].dt.date
```

/tmp/ipython-input-577587950.py:2: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df['issue_d'] = pd.to_datetime(df['issue_d'], errors='coerce')
```

/tmp/ipython-input-577587950.py:3: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'],
errors='coerce')
```

```
[ ]: df['issue_d_year'] = df['issue_d'].apply(lambda x: int(x.strftime('%Y')) if pd.
↳notnull(x) else None)
df['issue_d_month'] = df['issue_d'].apply(lambda x: int(x.strftime('%m')) if pd.
↳notnull(x) else None)
df['issue_d_day'] = df['issue_d'].apply(lambda x: int(x.strftime('%d')) if pd.
↳notnull(x) else None)

df['earliest_cr_line_year'] = df['earliest_cr_line'].apply(lambda x: int(x.
↳strftime('%Y')) if pd.notnull(x) else None)
df['earliest_cr_line_month'] = df['earliest_cr_line'].apply(lambda x: int(x.
↳strftime('%m')) if pd.notnull(x) else None)
df['earliest_cr_line_day'] = df['earliest_cr_line'].apply(lambda x: int(x.
↳strftime('%d')) if pd.notnull(x) else None)
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 324859 entries, 0 to 396028
```

```
Data columns (total 34 columns):
```

#	Column	Non-Null Count	Dtype
0	loan_amnt	324859 non-null	float64
1	term	324859 non-null	object
2	int_rate	324859 non-null	float64
3	installment	324859 non-null	float64
4	grade	324859 non-null	object
5	sub_grade	324859 non-null	object
6	emp_title	321562 non-null	object
7	emp_length	324859 non-null	object
8	home_ownership	324859 non-null	object
9	annual_inc	324859 non-null	float64
10	verification_status	324859 non-null	object
11	issue_d	324859 non-null	object
12	loan_status	324859 non-null	object

```

13 purpose          324859 non-null object
14 title            323457 non-null object
15 dti              324859 non-null float64
16 earliest_cr_line 324859 non-null object
17 open_acc         324859 non-null float64
18 pub_rec          324859 non-null float64
19 revol_bal        324859 non-null float64
20 revol_util       324859 non-null float64
21 total_acc        324859 non-null float64
22 initial_list_status 324859 non-null object
23 application_type  324859 non-null object
24 mort_acc         324859 non-null float64
25 pub_rec_bankruptcies 324859 non-null float64
26 address          324859 non-null object
27 zip_code         324859 non-null object
28 issue_d_year     324859 non-null int64
29 issue_d_month    324859 non-null int64
30 issue_d_day      324859 non-null int64
31 earliest_cr_line_year 324859 non-null int64
32 earliest_cr_line_month 324859 non-null int64
33 earliest_cr_line_day 324859 non-null int64
dtypes: float64(12), int64(6), object(16)
memory usage: 86.7+ MB

```

```
[ ]: df.drop(columns = ['issue_d', 'earliest_cr_line', 'emp_title', 'address', 'application_type'],
             ↪inplace=True)
```

```
[ ]: df.shape
```

```
[ ]: (324859, 29)
```

```
[ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 324859 entries, 0 to 396028
Data columns (total 29 columns):
#   Column          Non-Null Count  Dtype
---  -
0   loan_amnt       324859 non-null float64
1   term            324859 non-null object
2   int_rate        324859 non-null float64
3   installment     324859 non-null float64
4   grade           324859 non-null object
5   sub_grade       324859 non-null object
6   emp_length      324859 non-null object
7   home_ownership  324859 non-null object
8   annual_inc      324859 non-null float64

```

```

9  verification_status      324859 non-null object
10 loan_status              324859 non-null object
11 purpose                  324859 non-null object
12 title                    323457 non-null object
13 dti                      324859 non-null float64
14 open_acc                 324859 non-null float64
15 pub_rec                  324859 non-null float64
16 revol_bal                324859 non-null float64
17 revol_util               324859 non-null float64
18 total_acc                324859 non-null float64
19 initial_list_status      324859 non-null object
20 mort_acc                 324859 non-null float64
21 pub_rec_bankruptcies    324859 non-null float64
22 zip_code                 324859 non-null object
23 issue_d_year             324859 non-null int64
24 issue_d_month            324859 non-null int64
25 issue_d_day              324859 non-null int64
26 earliest_cr_line_year    324859 non-null int64
27 earliest_cr_line_month   324859 non-null int64
28 earliest_cr_line_day     324859 non-null int64
dtypes: float64(12), int64(6), object(11)
memory usage: 74.4+ MB

```

```
[ ]: df['statecode'] = df['zip_code'].str.extract(r'^(([A-Z]{2}))')
df['zip'] = df['zip_code'].str.extract(r'(\d{5})$')
```

```
[ ]: df.drop(columns = ['zip_code'], inplace=True)
```

```
[ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 324859 entries, 0 to 396028
Data columns (total 30 columns):
#   Column              Non-Null Count  Dtype
---  -
0   loan_amnt           324859 non-null float64
1   term                324859 non-null object
2   int_rate            324859 non-null float64
3   installment         324859 non-null float64
4   grade               324859 non-null object
5   sub_grade           324859 non-null object
6   emp_length          324859 non-null object
7   home_ownership      324859 non-null object
8   annual_inc          324859 non-null float64
9   verification_status 324859 non-null object
10  loan_status         324859 non-null object
11  purpose              324859 non-null object
12  title                323457 non-null object

```

```

13 dti 324859 non-null float64
14 open_acc 324859 non-null float64
15 pub_rec 324859 non-null float64
16 revol_bal 324859 non-null float64
17 revol_util 324859 non-null float64
18 total_acc 324859 non-null float64
19 initial_list_status 324859 non-null object
20 mort_acc 324859 non-null float64
21 pub_rec_bankruptcies 324859 non-null float64
22 issue_d_year 324859 non-null int64
23 issue_d_month 324859 non-null int64
24 issue_d_day 324859 non-null int64
25 earliest_cr_line_year 324859 non-null int64
26 earliest_cr_line_month 324859 non-null int64
27 earliest_cr_line_day 324859 non-null int64
28 statecode 324859 non-null object
29 zip 324859 non-null object
dtypes: float64(12), int64(6), object(12)
memory usage: 76.8+ MB

```

```
[ ]: df['zip'] = df['zip'].astype(int)
```

```
[ ]: sorted_unique_values = sorted(df['emp_length'].unique())
sorted_unique_values
```

```
[ ]: ['1 year',
      '10+ years',
      '2 years',
      '3 years',
      '4 years',
      '5 years',
      '6 years',
      '7 years',
      '8 years',
      '9 years',
      '< 1 year']
```

```
[ ]: #rename <1 year to 0-1 year, 10+ years to 10 years+
df['emp_length'] = df['emp_length'].replace(['< 1 year'], '0-1 year')
df['emp_length'] = df['emp_length'].replace(['10+ years'], '9 years+')

```

```
[ ]: sorted_unique_values1 = sorted(df['emp_length'].unique())
sorted_unique_values1
```

```
[ ]: ['0-1 year',
      '1 year',
      '2 years',
```

```
'3 years',
'4 years',
'5 years',
'6 years',
'7 years',
'8 years',
'9 years',
'9 years+']
```

```
[ ]: #save the dataset now
df.to_csv('loan_data_cleaned.csv', index=False)
```

```
[ ]: #drop statecode, title columns
#purpose and title are almost similar classes
df.drop(columns = ['statecode', 'title', 'sub_grade', 'dti', 'pub_rec'],
        inplace=True)
```

```
[ ]: df.shape
```

```
[ ]: (324859, 25)
```

##One Hot Encoding

```
[ ]: from sklearn.preprocessing import OneHotEncoder
import pandas as pd

# Select the columns you want to encode
cols_to_encode = ['home_ownership', 'purpose']

# Initialize the encoder
ohe = OneHotEncoder(sparse_output = False, handle_unknown='ignore')

# Fit and transform
encoded = ohe.fit_transform(df[cols_to_encode])

# Convert to DataFrame and preserve index
encoded_df = pd.DataFrame(encoded, columns=ohe.
    get_feature_names_out(cols_to_encode), index=df.index)

# Concatenate back with original dataframe (drop the encoded columns first)
df_encoded = pd.concat([df.drop(cols_to_encode, axis=1), encoded_df], axis=1)
```

```
[ ]: df_encoded.shape
```

```
[ ]: (324859, 40)
```

##Label Encoding


```
[ ]: #Label encoding for categorical features
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df_encoded['term'] = le.fit_transform(df_encoded['term'])
df_encoded['emp_length'] = le.fit_transform(df_encoded['emp_length'])
df_encoded['loan_status'] = le.fit_transform(df_encoded['loan_status'])
df_encoded['grade'] = le.fit_transform(df_encoded['grade'])
df_encoded['verification_status'] = le.
    ↪fit_transform(df_encoded['verification_status'])
df_encoded['initial_list_status'] = le.
    ↪fit_transform(df_encoded['initial_list_status'])
```

```
[ ]: df_encoded.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 324859 entries, 0 to 396028
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                             324859 non-null  float64
1   term                                  324859 non-null  int64
2   int_rate                              324859 non-null  float64
3   installment                           324859 non-null  float64
4   grade                                  324859 non-null  int64
5   emp_length                             324859 non-null  int64
6   annual_inc                             324859 non-null  float64
7   verification_status                   324859 non-null  int64
8   loan_status                           324859 non-null  int64
9   open_acc                              324859 non-null  float64
10  revol_bal                              324859 non-null  float64
11  revol_util                             324859 non-null  float64
12  total_acc                              324859 non-null  float64
13  initial_list_status                   324859 non-null  int64
14  mort_acc                              324859 non-null  float64
15  pub_rec_bankruptcies                  324859 non-null  float64
16  issue_d_year                          324859 non-null  int64
17  issue_d_month                         324859 non-null  int64
18  issue_d_day                           324859 non-null  int64
19  earliest_cr_line_year                  324859 non-null  int64
20  earliest_cr_line_month                 324859 non-null  int64
21  earliest_cr_line_day                   324859 non-null  int64
22  zip                                    324859 non-null  int64
23  home_ownership_ANY                     324859 non-null  float64
24  home_ownership_MORTGAGE                 324859 non-null  float64
25  home_ownership_OWN                     324859 non-null  float64
26  home_ownership_RENT                     324859 non-null  float64
27  purpose_car                            324859 non-null  float64
```


28	purpose_credit_card	324859	non-null	float64
29	purpose_debt_consolidation	324859	non-null	float64
30	purpose_home_improvement	324859	non-null	float64
31	purpose_house	324859	non-null	float64
32	purpose_major_purchase	324859	non-null	float64
33	purpose_medical	324859	non-null	float64
34	purpose_moving	324859	non-null	float64
35	purpose_other	324859	non-null	float64
36	purpose_renewable_energy	324859	non-null	float64
37	purpose_small_business	324859	non-null	float64
38	purpose_vacation	324859	non-null	float64
39	purpose_wedding	324859	non-null	float64

dtypes: float64(27), int64(13)

memory usage: 101.6 MB

```
[ ]: #Correlation
corr1 = df_encoded.corr()
plt.figure(figsize=(30,10))
sns.heatmap(corr1, annot=True)
```

[]: <Axes: >



#ML Modelling

```
[ ]: #X and Y value


# Define the target variable (dependent variable)
y = df_encoded['loan_status']

# Define the feature set (independent variables)
```

```
X = df_encoded.drop('loan_status', axis=1)
```

```
[ ]: #Pie chart for loan_status before SMOTE
loan_status_counts = df_encoded['loan_status'].value_counts()

# Plot
plt.figure(figsize=(8, 8))
labels = loan_status_counts
plt.pie(loan_status_counts, labels=loan_status_counts, autopct='%1.1f%%',
        ↪startangle=140)
plt.title('Distribution of Loan Status')
plt.axis('equal') # Equal aspect ratio ensures the pie is a circle.
plt.show()
```




##SMOTE to handle imbalance data

```
[ ]: #SMOTE to handle imbalance data in loan status feature
from imblearn.under_sampling import RandomUnderSampler
undersampler = RandomUnderSampler(random_state=42)
# Drop columns before undersampling
X = X.drop(['earliest_cr_line_day', 'issue_d_day'], axis=1)
X_res, y_res = undersampler.fit_resample(X, y)
```

```
[ ]: #Pie chart for loan_status after SMOTE
loan_status_counts1 = y_res.value_counts()

# Plot
plt.figure(figsize=(8, 8))
plt.pie(loan_status_counts1, labels=loan_status_counts1, autopct='%1.1f%%',
↪startangle=140)
plt.title('Distribution of Loan Status')
plt.axis('equal') # Equal aspect ratio ensures the pie is a circle.
plt.show()
```

Distribution of Loan Status



```
[ ]: X_res.shape
```

```
[ ]: (126766, 37)
```

```
##Machine Learning Algorithms
```

```
[ ]: #Test Train data split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.
↪2, random_state=42)
```

```
###Logistic Regression
```

```
[ ]: #Logistic Classifier model
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(X_train, y_train)
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[ ]: LogisticRegression()
```

```
[ ]: #Predict and find accuracy
from sklearn.metrics import accuracy_score
y_pred_lr = lr.predict(X_test)
accuracy = accuracy_score(y_test, y_pred_lr)
print("Accuracy:", accuracy)
```

Accuracy: 0.6684152401987852

```
[ ]: #hyperparameter tuning for logistic regression
..from sklearn.model_selection import GridSearchCV
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2'],
    'solver': ['liblinear', 'saga']
}
grid_search = GridSearchCV(estimator=lr, param_grid=param_grid, cv=5,
    ↪scoring='accuracy')
grid_search.fit(X_train_cv, y_train_cv)
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
best_params
```

```
File "/tmp/ipython-input-3370312101.py", line 2
    ..from sklearn.model_selection import GridSearchCV
    ^
```

SyntaxError: invalid syntax

C=10, penalty='l1', solver='liblinear'

```
[ ]: best_model = LogisticRegression(C=10, penalty='l1', solver='liblinear')
best_model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(C=10, penalty='l1', solver='liblinear')
```

```
[ ]: from sklearn.metrics import classification_report, confusion_matrix,
      accuracy_score

y_pred_lr_best = best_model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred_lr_best))
print(confusion_matrix(y_test, y_pred_lr_best))
print(classification_report(y_test, y_pred_lr_best))
```

Accuracy: 0.7085666955904394

[[8836 3947]

[3442 9129]]

	precision	recall	f1-score	support
0	0.72	0.69	0.71	12783
1	0.70	0.73	0.71	12571
accuracy			0.71	25354
macro avg	0.71	0.71	0.71	25354
weighted avg	0.71	0.71	0.71	25354

Feature Importance for Logistic Regression

```
[ ]: ### 2. Feature Selection using Logistic Regression (L1 penalty) ###
from sklearn.feature_selection import SelectFromModel
import pandas as pd
log_reg = LogisticRegression(penalty='l1', solver='liblinear', C=1.0,
                             random_state=42, max_iter=1000)
selector = SelectFromModel(log_reg)
selector.fit(X_train, y_train)

selected_features = X_train.columns[selector.get_support()]
print("\nSelected Features from Logistic Regression (L1):")
print(selected_features.tolist())
```

Selected Features from Logistic Regression (L1):

```
['loan_amnt', 'term', 'int_rate', 'installment', 'grade', 'emp_length',
'verification_status', 'open_acc', 'revol_util', 'total_acc',
'initial_list_status', 'mort_acc', 'pub_rec_bankruptcies', 'issue_d_year',
'issue_d_month', 'earliest_cr_line_year', 'earliest_cr_line_month', 'zip',
'home_ownership_MORTGAGE', 'home_ownership_OWN', 'home_ownership_RENT',
```

```
'purpose_car', 'purpose_credit_card', 'purpose_home_improvement',
'purpose_house', 'purpose_major_purchase', 'purpose_medical', 'purpose_moving',
'purpose_other', 'purpose_renewable_energy', 'purpose_small_business',
'purpose_vacation', 'purpose_wedding']
```

```
[ ]: #Top features from Logistic Regression (L1) to new df
# Include 'loan_status' in the selection
# Sort the selected_features index and take the top 15
top_lr_features = selected_features.sort_values(ascending=False).tolist()[:15]
new_lr_df = df_encoded[top_lr_features + ['loan_status']]
```

```
[ ]: #Split to X and y, test and train data and run LR model
y_new_df = new_lr_df['loan_status']
X_new_df = new_lr_df.drop('loan_status', axis=1)

X_new_train, X_new_test, y_new_train, y_new_test = train_test_split(X_new_df,
    ↪y_new_df, test_size=0.2, random_state=42)

#LR model
lr_new = LogisticRegression()
lr_new.fit(X_new_train, y_new_train)

#predict accuracy
y_pred_lr_new = lr_new.predict(X_new_test)
accuracy = accuracy_score(y_new_test, y_pred_lr_new)
print("Accuracy:", accuracy)
print(classification_report(y_new_test, y_pred_lr_new))
```

Accuracy: 0.836175583328203

	precision	recall	f1-score	support
0	0.73	0.27	0.39	12793
1	0.84	0.98	0.91	52179
accuracy			0.84	64972
macro avg	0.79	0.62	0.65	64972
weighted avg	0.82	0.84	0.80	64972

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:


<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression


```
n_iter_i = _check_optimize_result(
```

```
[ ]: #Confusion Matrix and AUC-ROC curve code
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, roc_curve, auc, \
    ConfusionMatrixDisplay
# Confusion Matrix
cm = confusion_matrix(y_new_test, y_pred_lr_new)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix")
plt.show()
```




```
[ ]: # 6. ROC Curve and AUC
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Calculate predicted probabilities for the positive class
y_proba = lr_new.predict_proba(X_new_test)[: , 1]
```

```
fpr, tpr, thresholds = roc_curve(y_new_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' %
        roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.grid(True)
plt.show()
```



```
[ ]: #save lr model
      #import joblib
      #joblib.dump(lr_new, 'lr_model.pkl')
```

###Gradient Boosting Method

```
[ ]: #GBM machine learning model
from sklearn.ensemble import GradientBoostingClassifier
gbm = GradientBoostingClassifier()
gbm.fit(X_train, y_train)
```

```
[ ]: GradientBoostingClassifier()
```

```
[ ]: #Predict and find accuracy
from sklearn.metrics import accuracy_score
y_pred_gbm = gbm.predict(X_test)
accuracy = accuracy_score(y_test, y_pred_gbm)
print("Accuracy:", accuracy)
```

Accuracy: 0.8063027530172754

```
[ ]: #hyperparameter tuning for GBM model
from sklearn.model_selection import GridSearchCV
param_grid = {
    'n_estimators': [100, 150],
    'learning_rate': [0.05, 0.1],
    'max_depth': [3, 4]
}
grid_search = GridSearchCV(estimator=gbm, param_grid=param_grid, cv=5,
    ↪scoring='accuracy')
grid_search.fit(X_train_cv, y_train_cv)
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
best_params
```

learning_rate=0.1, max_depth=4, n_estimators=150

```
[ ]: #GBM machine learning model
from sklearn.ensemble import GradientBoostingClassifier
gbm_h = GradientBoostingClassifier(learning_rate=0.1, max_depth=4,
    ↪n_estimators=150)
gbm_h.fit(X_train, y_train)
```

```
[ ]: GradientBoostingClassifier(max_depth=4, n_estimators=150)
```

```
[ ]: #Predict and find accuracy
from sklearn.metrics import accuracy_score
y_pred_gbm_h = gbm_h.predict(X_test)
accuracy = accuracy_score(y_test, y_pred_gbm_h)
print("Accuracy:", accuracy)
```

Accuracy: 0.8082748284294391

```
[ ]: #Classification matrix or F1-score
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_gbm))
```

	precision	recall	f1-score	support
0	0.80	0.82	0.81	12783
1	0.81	0.79	0.80	12571
accuracy			0.81	25354
macro avg	0.81	0.81	0.81	25354
weighted avg	0.81	0.81	0.81	25354

Feature Importance for GBM

```
[ ]: # 1. Feature importance using Gradient Boosting Machine (GBM)
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.feature_selection import SelectFromModel
import pandas as pd

# Assuming 'gbm' model is already trained
# gbm = GradientBoostingClassifier()
# gbm.fit(X_train, y_train)

feature_importance_gbm = pd.Series(gbm.feature_importances_, index=X_train.
    ↪columns)
print("\nTop Features from GBM:")
print(feature_importance_gbm.sort_values(ascending=False).head(15))
```

Top Features from GBM:

zip	0.829751
grade	0.106077
int_rate	0.015608
term	0.014113
issue_d_year	0.010837
annual_inc	0.007725
issue_d_month	0.003024
mort_acc	0.002483
installment	0.001929
revol_util	0.001684
home_ownership_RENT	0.001611
loan_amnt	0.001484
open_acc	0.001239
home_ownership_MORTGAGE	0.001130
verification_status	0.000569

dtype: float64

```
[ ]: #Top features from GBM to new df
# Include 'loan_status' in the selection
new_df = df_encoded[feature_importance_gbm.sort_values(ascending=False).
↳head(15).index.tolist() + ['loan_status']]
```

```
[ ]: new_df.head()
```

```
[ ]:      zip  grade  int_rate  term  issue_d_year  annual_inc  issue_d_month  \
0  22690      1    11.44     0        2015    117000.0           1
1   5113      1    11.99     0        2015     65000.0           1
2   5113      1    10.49     0        2015    43057.0           1
3    813      0     6.49     0        2014    54000.0          11
4  11650      2    17.27     1        2013    55000.0           4
```

```
      mort_acc  installment  revol_util  home_ownership_RENT  loan_amnt  \
0         0.0        329.48        41.8             1.0    10000.0
1         3.0        265.68        53.3             0.0     8000.0
2         0.0        506.97        92.2             1.0    15600.0
3         0.0        220.65        21.5             1.0     7200.0
4         1.0        609.33        69.8             0.0    24375.0
```

```
      open_acc  home_ownership_MORTGAGE  verification_status  loan_status
0         16.0                0.0                0           1
1         17.0                1.0                0           1
2         13.0                0.0                1           1
3          6.0                0.0                0           1
4         13.0                1.0                2           0
```

```
[ ]: #Split to X and y, test and train data and run GBM model
y_new_df = new_df['loan_status']
X_new_df = new_df.drop('loan_status', axis=1)

X_new_train, X_new_test, y_new_train, y_new_test = train_test_split(X_new_df,
↳y_new_df, test_size=0.2, random_state=42)

#GBM model
gbm_new = GradientBoostingClassifier()
gbm_new.fit(X_new_train, y_new_train)


#predict accuracy
y_pred_gbm_new = gbm_new.predict(X_new_test)
accuracy = accuracy_score(y_new_test, y_pred_gbm_new)
print("Accuracy:", accuracy)
```

Accuracy: 0.8913070245644278

```
[ ]: print(classification_report(y_new_test, y_pred_gbm_new))
```

	precision	recall	f1-score	support
0	0.96	0.47	0.63	12793
1	0.88	1.00	0.94	52179
accuracy			0.89	64972
macro avg	0.92	0.73	0.78	64972
weighted avg	0.90	0.89	0.88	64972

```
[ ]: # Confusion Matrix
cm = confusion_matrix(y_new_test, y_pred_gbm_new)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix")
plt.show()
```




```
[ ]: # 6. ROC Curve and AUC
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Calculate predicted probabilities for the positive class
y_proba = gbm_new.predict_proba(X_new_test)[: , 1]

fpr, tpr, thresholds = roc_curve(y_new_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' %
        roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.grid(True)
plt.show()
```



```
[ ]: #save gbm model
      #import joblib
      #joblib.dump(gbm_new, 'gbm_model.pkl')
```

###Random Forest

```
[ ]: #Random Forest model
      from sklearn.ensemble import RandomForestClassifier
      rf = RandomForestClassifier()
      rf.fit(X_train, y_train)
```

```
[ ]: RandomForestClassifier()
```

```
[ ]: #Predict and find accuracy
      from sklearn.metrics import accuracy_score
      y_pred_rf = rf.predict(X_test)
      accuracy = accuracy_score(y_test, y_pred_rf)
      print("Accuracy:", accuracy)
```

Accuracy: 0.7964818174646998

```
[ ]: #Hyperparameter tuning for Random Forest
      from sklearn.model_selection import GridSearchCV
      param_grid = {
          'n_estimators': [100, 150, 200],
          'max_depth': [None, 10, 20],
          'min_samples_split': [2, 5, 10],
          'min_samples_leaf': [1, 2, 4]
      }
      grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5,
          ↪scoring='accuracy')
      grid_search.fit(X_train_cv, y_train_cv)
      best_params = grid_search.best_params_
      best_model = grid_search.best_estimator_
      best_params
```

{'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 10, 'n_estimators': 150}

```
[ ]: rf_h = RandomForestClassifier(max_depth=None,
      min_samples_leaf=1,
      min_samples_split=10,
      n_estimators=150)
      rf_h.fit(X_train, y_train)
```

```
[ ]: RandomForestClassifier(min_samples_split=10, n_estimators=150)
```



```
[ ]: #Predict and find accuracy
from sklearn.metrics import accuracy_score
y_pred_rf = rf_h.predict(X_test)
accuracy = accuracy_score(y_test, y_pred_rf)
print("Accuracy:", accuracy)
```

Accuracy: 0.7966790250059163

Feature Importance for Random Forest

```
[ ]: # 1. Feature importance using Random Forest
from sklearn.feature_selection import SelectKBest, f_classif

feature_importances = pd.Series(rf.feature_importances_, index=X_train.columns)
important_features_rf = feature_importances.sort_values(ascending=False)

print("\nTop Features from Random Forest:")
print(important_features_rf.head(10))
```

```
Top Features from Random Forest:
zip                0.307047
int_rate           0.069580
revol_util         0.053619
annual_inc         0.052212
revol_bal          0.051842
installment        0.049085
grade              0.046517
total_acc          0.041708
loan_amnt          0.040311
earliest_cr_line_year 0.039036
dtype: float64
```

```
[ ]: top_rf_features = important_features_rf.sort_values(ascending=False).head(15).
    ↪ index.tolist()
new_rf_df = df_encoded[top_rf_features + ['loan_status']]
```

```
[ ]: #Split to X and y, test and train data and run RF model
y_new_df = new_rf_df['loan_status']
X_new_df = new_rf_df.drop('loan_status', axis=1)

X_new_train, X_new_test, y_new_train, y_new_test = train_test_split(X_new_df,
    ↪ y_new_df, test_size=0.2, random_state=42)

#LR model
rf_new = RandomForestClassifier()
rf_new.fit(X_new_train, y_new_train)
```

```

#predict accuracy
y_pred_rf_new = rf_new.predict(X_new_test)
accuracy = accuracy_score(y_new_test, y_pred_rf_new)
print("Accuracy:", accuracy)
print(classification_report(y_new_test, y_pred_rf_new))

```


Accuracy: 0.8903989410823124

	precision	recall	f1-score	support
0	0.94	0.47	0.63	12793
1	0.88	0.99	0.94	52179
accuracy			0.89	64972
macro avg	0.91	0.73	0.78	64972
weighted avg	0.90	0.89	0.88	64972

```

[ ]: #Confusion Matrix and AUC-ROC curve code
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, roc_curve, auc, ConfusionMatrixDisplay
# Confusion Matrix
cm = confusion_matrix(y_new_test, y_pred_rf_new)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix")
plt.show()

```




```
[ ]: # 6. ROC Curve and AUC
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Calculate predicted probabilities for the positive class
y_proba = rf_new.predict_proba(X_new_test)[: , 1]

fpr, tpr, thresholds = roc_curve(y_new_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.grid(True)
plt.show()
```



```
[ ]: #save rf model
      #import joblib
      #joblib.dump(rf, 'rf_model.pkl')
```

###Support Vector Machine

```
[ ]: from sklearn.svm import LinearSVC
      from sklearn.calibration import CalibratedClassifierCV

      # LinearSVC does NOT support predict_proba directly
      base_model = LinearSVC(random_state=42, max_iter=10000)
      svm = CalibratedClassifierCV(base_model)

      svm.fit(X_train, y_train)
```

```
[ ]: CalibratedClassifierCV(estimator=LinearSVC(max_iter=10000, random_state=42))
```

```
[ ]: #Predict and find accuracy
      from sklearn.metrics import accuracy_score
```

```

y_pred_svm = svm.predict(X_test)
accuracy = accuracy_score(y_test, y_pred_svm)
print("Accuracy:", accuracy)

```

Accuracy: 0.7016644316478662

```

[ ]: #Hyperparameter tuning for SVM
from sklearn.model_selection import GridSearchCV
param_grid = {
    'C': [0.1, 1, 10],
    'kernel': ['linear', 'rbf'],
    'gamma': ['scale', 'auto']
}
grid_search = GridSearchCV(estimator=svm, param_grid=param_grid, cv=5,
    ↪scoring='accuracy')
grid_search.fit(X_train_cv, y_train_cv)
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
best_params

```

Feature Importance for SVM

```

[ ]: from sklearn.inspection import permutation_importance

# Use a smaller sample from X_test if dataset is large
sample_size = 500 # Adjust as needed
X_test_sample = X_test[:sample_size]
y_test_sample = y_test[:sample_size]

result = permutation_importance(
    svm,
    X_test_sample,
    y_test_sample,
    n_repeats=5,
    random_state=42,
    n_jobs=-1
)

# If using NumPy array instead of DataFrame
if isinstance(X_test_sample, pd.DataFrame):
    feature_names = X_test_sample.columns
else:
    feature_names = [f"Feature_{i}" for i in range(X_test_sample.shape[1])]

# Organize and sort importances
perm_importance_svm = pd.Series(result.importances_mean, index=feature_names)
perm_importance_svm = perm_importance_svm.sort_values(ascending=False)

```

```
# Print top features
print("\nTop Feature Importances from Permutation Importance (SVM):")
print(perm_importance_svm.head(10))
```

Top Feature Importances from Permutation Importance (SVM):

```
zip                0.1296
installment        0.0272
int_rate           0.0184
loan_amnt          0.0160
annual_inc         0.0116
open_acc           0.0084
total_acc          0.0080
issue_d_month      0.0064
emp_length         0.0060
mort_acc           0.0056
dtype: float64
```

```
[ ]: top_svm_features = perm_importance_svm.sort_values(ascending=False).head(15).
      ↪index.tolist()
new_svm_df = df_encoded[top_svm_features + ['loan_status']]
```

```
[ ]: #Split to X and y, test and train data and run SVM model
y_new_df = new_svm_df['loan_status']
X_new_df = new_svm_df.drop('loan_status', axis=1)

X_new_train, X_new_test, y_new_train, y_new_test = train_test_split(X_new_df,
      ↪y_new_df, test_size=0.2, random_state=42)

#SVM model
svm_new = LinearSVC(random_state=42, max_iter=10000)
svm_new.fit(X_new_train, y_new_train)


#predict accuracy
y_pred_svm_new = svm_new.predict(X_new_test)
accuracy = accuracy_score(y_new_test, y_pred_svm_new)
print("Accuracy:", accuracy)
print(classification_report(y_new_test, y_pred_svm_new))
```

Accuracy: 0.8399156559748815

	precision	recall	f1-score	support
0	0.93	0.20	0.33	12793
1	0.84	1.00	0.91	52179
accuracy			0.84	64972

macro avg	0.88	0.60	0.62	64972
weighted avg	0.85	0.84	0.80	64972

```
[ ]: #Confusion Matrix and AUC-ROC curve code
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, roc_curve, auc, ConfusionMatrixDisplay
# Confusion Matrix
cm = confusion_matrix(y_new_test, y_pred_svm_new)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix")
plt.show()
```



```
[ ]: # 6. ROC Curve and AUC
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Calculate predicted probabilities for the positive class
```


```

y_proba = svm.predict_proba(X_test)[: , 1]

fpr, tpr, thresholds = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.grid(True)
plt.show()

```



```

[ ]: #save svm model
      #import joblib
      #joblib.dump(svm, 'svm_model.pkl')

```


#ANN (Deep Learning)

```
[ ]: #Run ANN for the current dataset
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

#Artificial Neural network model
model = Sequential()
model.add(Dense(64, activation='relu', input_dim=X_new_train.shape[1]))
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

#predict ANN accuracy
model.compile(optimizer='adam', loss='binary_crossentropy',
              metrics=['accuracy'])
model.fit(X_new_train, y_new_train, epochs=20, batch_size=32,
          validation_split=0.2)
loss, accuracy = model.evaluate(X_new_test, y_new_test)
print(f"Test Loss: {loss:.4f}, Test Accuracy: {accuracy:.4f}")
```

Epoch 1/20

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

6498/6498 18s 3ms/step -
accuracy: 0.7302 - loss: 66.3755 - val_accuracy: 0.8278 - val_loss: 8.9592

Epoch 2/20

6498/6498 22s 3ms/step -
accuracy: 0.7620 - loss: 13.9426 - val_accuracy: 0.8104 - val_loss: 7.6403

Epoch 3/20

6498/6498 20s 3ms/step -
accuracy: 0.7820 - loss: 5.7421 - val_accuracy: 0.8221 - val_loss: 7.3997

Epoch 4/20

6498/6498 23s 3ms/step -
accuracy: 0.7950 - loss: 1.8184 - val_accuracy: 0.8843 - val_loss: 0.4140

Epoch 5/20

6498/6498 18s 3ms/step -
accuracy: 0.8438 - loss: 0.5632 - val_accuracy: 0.8768 - val_loss: 0.3450

Epoch 6/20

6498/6498 17s 3ms/step -
accuracy: 0.8723 - loss: 0.3405 - val_accuracy: 0.8149 - val_loss: 0.4267

Epoch 7/20

```

6498/6498          17s 3ms/step -
accuracy: 0.8240 - loss: 0.4129 - val_accuracy: 0.8076 - val_loss: 0.4417
Epoch 8/20
6498/6498          21s 3ms/step -
accuracy: 0.8287 - loss: 0.4048 - val_accuracy: 0.8193 - val_loss: 0.4161
Epoch 9/20
6498/6498          17s 3ms/step -
accuracy: 0.8163 - loss: 0.4211 - val_accuracy: 0.8040 - val_loss: 0.4769
Epoch 10/20
6498/6498          28s 4ms/step -
accuracy: 0.8174 - loss: 0.4212 - val_accuracy: 0.8215 - val_loss: 0.4125
Epoch 11/20
6498/6498          35s 3ms/step -
accuracy: 0.8220 - loss: 0.4141 - val_accuracy: 0.8039 - val_loss: 0.4560
Epoch 12/20
6498/6498          18s 3ms/step -
accuracy: 0.8054 - loss: 0.4378 - val_accuracy: 0.8039 - val_loss: 0.4361
Epoch 13/20
6498/6498          20s 3ms/step -
accuracy: 0.8046 - loss: 0.4361 - val_accuracy: 0.8039 - val_loss: 0.4361
Epoch 14/20
6498/6498          21s 3ms/step -
accuracy: 0.8052 - loss: 0.4356 - val_accuracy: 0.8039 - val_loss: 0.4361
Epoch 15/20
6498/6498          17s 3ms/step -
accuracy: 0.8046 - loss: 0.4369 - val_accuracy: 0.8039 - val_loss: 0.4360
Epoch 16/20
6498/6498          17s 3ms/step -
accuracy: 0.8051 - loss: 0.4358 - val_accuracy: 0.8039 - val_loss: 0.4360
Epoch 17/20
6498/6498          23s 3ms/step -
accuracy: 0.8067 - loss: 0.4336 - val_accuracy: 0.8039 - val_loss: 0.4360
Epoch 18/20
6498/6498          16s 3ms/step -
accuracy: 0.8060 - loss: 0.4349 - val_accuracy: 0.8039 - val_loss: 0.4360
Epoch 19/20
6498/6498          22s 3ms/step -
accuracy: 0.8076 - loss: 0.4326 - val_accuracy: 0.8039 - val_loss: 0.4360
Epoch 20/20
6498/6498          17s 3ms/step -
accuracy: 0.8059 - loss: 0.4349 - val_accuracy: 0.8039 - val_loss: 0.4360
2031/2031          3s 2ms/step -
accuracy: 0.8032 - loss: 0.4375
Test Loss: 0.4385, Test Accuracy: 0.8031

```

```

[ ]: # 1. Scaling
      scaler = StandardScaler()

```

```

X_train_scaled = scaler.fit_transform(X_new_train)
X_test_scaled = scaler.transform(X_new_test)

# 2. Model
model = Sequential()
model.add(Dense(128, activation='relu', input_dim=X_new_train.shape[1]))
model.add(Dense(64, activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

# 3. Compile
model.compile(optimizer='adam', loss='binary_crossentropy',
              metrics=['accuracy'])

# 4. Optional: class weights
from sklearn.utils import class_weight
class_weights = class_weight.compute_class_weight(class_weight='balanced',
                                                  classes=np.unique(y_train), y=y_new_train)
class_weight_dict = dict(enumerate(class_weights))

# 5. Train
model.fit(X_train_scaled, y_new_train, epochs=50, batch_size=32,
          validation_split=0.2, class_weight=class_weight_dict)

# 6. Evaluate
loss, accuracy = model.evaluate(X_test_scaled, y_new_test)
print(f"Test Loss: {loss:.4f}, Test Accuracy: {accuracy:.4f}")

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87:
 UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
 using Sequential models, prefer using an `Input(shape)` object as the first
 layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```

Epoch 1/50
6498/6498          21s 3ms/step -
accuracy: 0.7400 - loss: 0.5232 - val_accuracy: 0.7733 - val_loss: 0.4225
Epoch 2/50
6498/6498          22s 3ms/step -
accuracy: 0.8042 - loss: 0.3907 - val_accuracy: 0.8208 - val_loss: 0.3663
Epoch 3/50
6498/6498          23s 3ms/step -
accuracy: 0.8055 - loss: 0.3888 - val_accuracy: 0.8115 - val_loss: 0.3709
Epoch 4/50
6498/6498          41s 3ms/step -
accuracy: 0.8054 - loss: 0.3855 - val_accuracy: 0.8056 - val_loss: 0.3822
Epoch 5/50
6498/6498          38s 3ms/step -

```

accuracy: 0.8098 - loss: 0.3820 - val_accuracy: 0.7939 - val_loss: 0.3835
 Epoch 6/50
 6498/6498 22s 3ms/step -
 accuracy: 0.8052 - loss: 0.3820 - val_accuracy: 0.8017 - val_loss: 0.3782
 Epoch 7/50
 6498/6498 20s 3ms/step -
 accuracy: 0.8082 - loss: 0.3781 - val_accuracy: 0.8095 - val_loss: 0.3662
 Epoch 8/50
 6498/6498 21s 3ms/step -
 accuracy: 0.8094 - loss: 0.3699 - val_accuracy: 0.8037 - val_loss: 0.3617
 Epoch 9/50
 6498/6498 19s 3ms/step -
 accuracy: 0.8088 - loss: 0.3586 - val_accuracy: 0.8168 - val_loss: 0.3417
 Epoch 10/50
 6498/6498 19s 3ms/step -
 accuracy: 0.8038 - loss: 0.3559 - val_accuracy: 0.8266 - val_loss: 0.3318
 Epoch 11/50
 6498/6498 21s 3ms/step -
 accuracy: 0.8072 - loss: 0.3524 - val_accuracy: 0.8111 - val_loss: 0.3434
 Epoch 12/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8046 - loss: 0.3499 - val_accuracy: 0.8117 - val_loss: 0.3359
 Epoch 13/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8066 - loss: 0.3476 - val_accuracy: 0.8058 - val_loss: 0.3435
 Epoch 14/50
 6498/6498 22s 3ms/step -
 accuracy: 0.8064 - loss: 0.3478 - val_accuracy: 0.8132 - val_loss: 0.3380
 Epoch 15/50
 6498/6498 19s 3ms/step -
 accuracy: 0.8040 - loss: 0.3471 - val_accuracy: 0.8001 - val_loss: 0.3538
 Epoch 16/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8065 - loss: 0.3487 - val_accuracy: 0.8032 - val_loss: 0.3474
 Epoch 17/50
 6498/6498 21s 3ms/step -
 accuracy: 0.8084 - loss: 0.3482 - val_accuracy: 0.7980 - val_loss: 0.3525
 Epoch 18/50
 6498/6498 41s 3ms/step -
 accuracy: 0.8043 - loss: 0.3462 - val_accuracy: 0.8045 - val_loss: 0.3593
 Epoch 19/50
 6498/6498 40s 3ms/step -
 accuracy: 0.8100 - loss: 0.3473 - val_accuracy: 0.7995 - val_loss: 0.3538
 Epoch 20/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8083 - loss: 0.3467 - val_accuracy: 0.7967 - val_loss: 0.3541
 Epoch 21/50
 6498/6498 21s 3ms/step -

accuracy: 0.8064 - loss: 0.3454 - val_accuracy: 0.8045 - val_loss: 0.3470
 Epoch 22/50
 6498/6498 19s 3ms/step -
 accuracy: 0.8043 - loss: 0.3468 - val_accuracy: 0.8070 - val_loss: 0.3532
 Epoch 23/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8083 - loss: 0.3461 - val_accuracy: 0.8198 - val_loss: 0.3375
 Epoch 24/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8081 - loss: 0.3455 - val_accuracy: 0.7904 - val_loss: 0.3620
 Epoch 25/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8069 - loss: 0.3457 - val_accuracy: 0.8150 - val_loss: 0.3413
 Epoch 26/50
 6498/6498 21s 3ms/step -
 accuracy: 0.8084 - loss: 0.3445 - val_accuracy: 0.7964 - val_loss: 0.3525
 Epoch 27/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8102 - loss: 0.3423 - val_accuracy: 0.8081 - val_loss: 0.3441
 Epoch 28/50
 6498/6498 21s 3ms/step -
 accuracy: 0.8092 - loss: 0.3433 - val_accuracy: 0.8072 - val_loss: 0.3496
 Epoch 29/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8077 - loss: 0.3445 - val_accuracy: 0.8109 - val_loss: 0.3443
 Epoch 30/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8099 - loss: 0.3418 - val_accuracy: 0.8115 - val_loss: 0.3432
 Epoch 31/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8083 - loss: 0.3435 - val_accuracy: 0.7989 - val_loss: 0.3535
 Epoch 32/50
 6498/6498 21s 3ms/step -
 accuracy: 0.8088 - loss: 0.3452 - val_accuracy: 0.7991 - val_loss: 0.3552
 Epoch 33/50
 6498/6498 18s 3ms/step -
 accuracy: 0.8107 - loss: 0.3436 - val_accuracy: 0.8173 - val_loss: 0.3320
 Epoch 34/50
 6498/6498 21s 3ms/step -
 accuracy: 0.8095 - loss: 0.3415 - val_accuracy: 0.8077 - val_loss: 0.3493
 Epoch 35/50
 6498/6498 20s 3ms/step -
 accuracy: 0.8117 - loss: 0.3422 - val_accuracy: 0.8079 - val_loss: 0.3436
 Epoch 36/50
 6498/6498 20s 3ms/step -
 accuracy: 0.8121 - loss: 0.3385 - val_accuracy: 0.8075 - val_loss: 0.3538
 Epoch 37/50
 6498/6498 20s 3ms/step -

```

accuracy: 0.8109 - loss: 0.3429 - val_accuracy: 0.7822 - val_loss: 0.3772
Epoch 38/50
6498/6498          18s 3ms/step -
accuracy: 0.8094 - loss: 0.3432 - val_accuracy: 0.8053 - val_loss: 0.3533
Epoch 39/50
6498/6498          23s 3ms/step -
accuracy: 0.8096 - loss: 0.3405 - val_accuracy: 0.7841 - val_loss: 0.3806
Epoch 40/50
6498/6498          17s 3ms/step -
accuracy: 0.8103 - loss: 0.3432 - val_accuracy: 0.7910 - val_loss: 0.3664
Epoch 41/50
6498/6498          18s 3ms/step -
accuracy: 0.8088 - loss: 0.3420 - val_accuracy: 0.7994 - val_loss: 0.3526
Epoch 42/50
6498/6498          18s 3ms/step -
accuracy: 0.8101 - loss: 0.3383 - val_accuracy: 0.8082 - val_loss: 0.3518
Epoch 43/50
6498/6498          19s 3ms/step -
accuracy: 0.8133 - loss: 0.3366 - val_accuracy: 0.8005 - val_loss: 0.3621
Epoch 44/50
6498/6498          18s 3ms/step -
accuracy: 0.8126 - loss: 0.3400 - val_accuracy: 0.8010 - val_loss: 0.3523
Epoch 45/50
6498/6498          18s 3ms/step -
accuracy: 0.8107 - loss: 0.3371 - val_accuracy: 0.8125 - val_loss: 0.3420
Epoch 46/50
6498/6498          18s 3ms/step -
accuracy: 0.8099 - loss: 0.3388 - val_accuracy: 0.8112 - val_loss: 0.3436
Epoch 47/50
6498/6498          20s 3ms/step -
accuracy: 0.8094 - loss: 0.3403 - val_accuracy: 0.7897 - val_loss: 0.3711
Epoch 48/50
6498/6498          18s 3ms/step -
accuracy: 0.8112 - loss: 0.3387 - val_accuracy: 0.8021 - val_loss: 0.3556
Epoch 49/50
6498/6498          20s 3ms/step -
accuracy: 0.8120 - loss: 0.3396 - val_accuracy: 0.8137 - val_loss: 0.3471
Epoch 50/50
6498/6498          19s 3ms/step -
accuracy: 0.8132 - loss: 0.3387 - val_accuracy: 0.7915 - val_loss: 0.3610
2031/2031          3s 2ms/step -
accuracy: 0.7916 - loss: 0.3622
Test Loss: 0.3630, Test Accuracy: 0.7922

```

#Best Model Evaluation

```
[ ]: import pandas as pd
import numpy as np
```

```

import os
import joblib
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score

# ----- Load Dataset -----
df_path = "/content/loan_data_cleaned.csv"
df = pd.read_csv(df_path)

# ----- Extract Year from Date Column -----
if 'earliest_cr_line' in df.columns:
    df['earliest_cr_line_year'] = pd.to_datetime(
        df['earliest_cr_line'], errors='coerce').dt.year

# ----- Define Feature Set -----
possible_features = [
    'zip', # optional
    'earliest_cr_line_year',
    'int_rate', 'revol_bal', 'revol_util',
    'annual_inc', 'installment', 'total_acc',
    'loan_amnt', 'open_acc'
]

# Use only existing features
features = [col for col in possible_features if col in df.columns]
missing = [col for col in possible_features if col not in df.columns]

print(f" Using features: {features}")
if missing:
    print(f" Skipping missing columns: {missing}")

# ----- Target Processing -----
target = 'loan_status'
if df[target].dtype == 'object':
    le = LabelEncoder()
    df[target] = le.fit_transform(df[target])

# ----- Feature / Target Split -----
X = df[features]
y = df[target]

# ----- Train-Test Split -----

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# ----- Define Pipeline -----
pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler()), # Optional but good for many ML models
    ('model', GradientBoostingClassifier())
])

# ----- Train Model -----
pipeline.fit(X_train, y_train)

# ----- Evaluate Model -----
y_pred = pipeline.predict(X_test)
print("\n Classification Report:\n", classification_report(y_test, y_pred))
print(f" Accuracy: {accuracy_score(y_test, y_pred)*100:.2f}%")

# ----- Save Pipeline & Features -----
os.makedirs("models", exist_ok=True)
joblib.dump(pipeline, "models/gbm_pipeline.pkl")
joblib.dump(features, "models/model_columns.pkl")
print("\n Pipeline saved as 'models/gbm_pipeline.pkl'")
print(" Feature columns saved as 'models/model_columns.pkl'")

```

Using features: ['zip', 'earliest_cr_line_year', 'int_rate', 'revol_bal', 'revol_util', 'annual_inc', 'installment', 'total_acc', 'loan_amnt', 'open_acc']

```

Classification Report:
              precision    recall  f1-score   support

     0       0.99         0.45         0.62       12793
     1       0.88         1.00         0.94       52179

 accuracy                   0.89       64972
 macro avg              0.94         0.72         0.78       64972
weighted avg              0.90         0.89         0.87       64972

```

Accuracy: 89.04%

Pipeline saved as 'models/gbm_pipeline.pkl'
Feature columns saved as 'models/model_columns.pkl'